

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

%matplotlib inline
plt.style.use('seaborn-poster')
```

```
In [3]: df = pd.read_csv('/Users/swllms/DAT-10-14-SW/class material/Unit3/Data/housing.csv')
```

```
In [4]: df.head()
```

Out[4]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LST
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5

```
In [5]: from sklearn.linear_model import LinearRegression
```

```
In [21]: lreg=LinearRegression()
```

```
In [12]: X = df.iloc[:, :13]
```

```
In [17]: y = df['PRICE']
```

```
In [20]: X
```

Out[20]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90

5	0.02985	0.0	2.18	0	0.458	6.430	58.7	6.0622	3	222	18.7	394.12
6	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	395.60
7	0.14455	12.5	7.87	0	0.524	6.172	96.1	5.9505	5	311	15.2	396.90
8	0.21124	12.5	7.87	0	0.524	5.631	100.0	6.0821	5	311	15.2	386.63
9	0.17004	12.5	7.87	0	0.524	6.004	85.9	6.5921	5	311	15.2	386.71
10	0.22489	12.5	7.87	0	0.524	6.377	94.3	6.3467	5	311	15.2	392.52
11	0.11747	12.5	7.87	0	0.524	6.009	82.9	6.2267	5	311	15.2	396.90
12	0.09378	12.5	7.87	0	0.524	5.889	39.0	5.4509	5	311	15.2	390.50
13	0.62976	0.0	8.14	0	0.538	5.949	61.8	4.7075	4	307	21.0	396.90
14	0.63796	0.0	8.14	0	0.538	6.096	84.5	4.4619	4	307	21.0	380.02
15	0.62739	0.0	8.14	0	0.538	5.834	56.5	4.4986	4	307	21.0	395.62
16	1.05393	0.0	8.14	0	0.538	5.935	29.3	4.4986	4	307	21.0	386.85
17	0.78420	0.0	8.14	0	0.538	5.990	81.7	4.2579	4	307	21.0	386.75
18	0.80271	0.0	8.14	0	0.538	5.456	36.6	3.7965	4	307	21.0	288.99
19	0.72580	0.0	8.14	0	0.538	5.727	69.5	3.7965	4	307	21.0	390.95
20	1.25179	0.0	8.14	0	0.538	5.570	98.1	3.7979	4	307	21.0	376.57
21	0.85204	0.0	8.14	0	0.538	5.965	89.2	4.0123	4	307	21.0	392.53
22	1.23247	0.0	8.14	0	0.538	6.142	91.7	3.9769	4	307	21.0	396.90
23	0.98843	0.0	8.14	0	0.538	5.813	100.0	4.0952	4	307	21.0	394.54
24	0.75026	0.0	8.14	0	0.538	5.924	94.1	4.3996	4	307	21.0	394.33
25	0.84054	0.0	8.14	0	0.538	5.599	85.7	4.4546	4	307	21.0	303.42
26	0.67191	0.0	8.14	0	0.538	5.813	90.3	4.6820	4	307	21.0	376.88
27	0.95577	0.0	8.14	0	0.538	6.047	88.8	4.4534	4	307	21.0	306.38
28	0.77299	0.0	8.14	0	0.538	6.495	94.4	4.4547	4	307	21.0	387.94
29	1.00245	0.0	8.14	0	0.538	6.674	87.3	4.2390	4	307	21.0	380.23
...
476	4.87141	0.0	18.10	0	0.614	6.484	93.6	2.3053	24	666	20.2	396.21
477	15.02340	0.0	18.10	0	0.614	5.304	97.3	2.1007	24	666	20.2	349.48
478	10.23300	0.0	18.10	0	0.614	6.185	96.7	2.1705	24	666	20.2	379.70
479	14.33370	0.0	18.10	0	0.614	6.229	88.0	1.9512	24	666	20.2	383.32
480	5.82401	0.0	18.10	0	0.532	6.242	64.7	3.4242	24	666	20.2	396.90
481	5.70818	0.0	18.10	0	0.532	6.750	74.9	3.3317	24	666	20.2	393.07

482	5.73116	0.0	18.10	0	0.532	7.061	77.0	3.4106	24	666	20.2	395.28
483	2.81838	0.0	18.10	0	0.532	5.762	40.3	4.0983	24	666	20.2	392.92
484	2.37857	0.0	18.10	0	0.583	5.871	41.9	3.7240	24	666	20.2	370.73
485	3.67367	0.0	18.10	0	0.583	6.312	51.9	3.9917	24	666	20.2	388.62
486	5.69175	0.0	18.10	0	0.583	6.114	79.8	3.5459	24	666	20.2	392.68
487	4.83567	0.0	18.10	0	0.583	5.905	53.2	3.1523	24	666	20.2	388.22
488	0.15086	0.0	27.74	0	0.609	5.454	92.7	1.8209	4	711	20.1	395.09
489	0.18337	0.0	27.74	0	0.609	5.414	98.3	1.7554	4	711	20.1	344.05
490	0.20746	0.0	27.74	0	0.609	5.093	98.0	1.8226	4	711	20.1	318.43
491	0.10574	0.0	27.74	0	0.609	5.983	98.8	1.8681	4	711	20.1	390.11
492	0.11132	0.0	27.74	0	0.609	5.983	83.5	2.1099	4	711	20.1	396.90
493	0.17331	0.0	9.69	0	0.585	5.707	54.0	2.3817	6	391	19.2	396.90
494	0.27957	0.0	9.69	0	0.585	5.926	42.6	2.3817	6	391	19.2	396.90
495	0.17899	0.0	9.69	0	0.585	5.670	28.8	2.7986	6	391	19.2	393.29
496	0.28960	0.0	9.69	0	0.585	5.390	72.9	2.7986	6	391	19.2	396.90
497	0.26838	0.0	9.69	0	0.585	5.794	70.6	2.8927	6	391	19.2	396.90
498	0.23912	0.0	9.69	0	0.585	6.019	65.3	2.4091	6	391	19.2	396.90
499	0.17783	0.0	9.69	0	0.585	5.569	73.5	2.3999	6	391	19.2	395.77
500	0.22438	0.0	9.69	0	0.585	6.027	79.7	2.4982	6	391	19.2	396.90
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396.90

506 rows × 13 columns

In [22]: `lreg.fit(X, y)`

Out[22]: `LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)`

In [35]: `lreg.score(X, y) #74%`

Out[35]: `0.7406426641094094`

In [23]: `lreg.coef_`

Out[23]: `array([-1.08011358e-01, 4.64204584e-02, 2.05586264e-02, 2.6867338
2e+00,
 -1.77666112e+01, 3.80986521e+00, 6.92224640e-04, -1.4755668
5e+00,
 3.06049479e-01, -1.23345939e-02, -9.52747232e-01, 9.3116832
7e-03,
 -5.24758378e-01])`

In [26]: `coeffs = pd.DataFrame({
 'Variable': X.columns,
 'Weight' : lreg.coef_
}).sort_values(by='Weight', ascending=False)`

In [27]: *coeffs #For every value of variable add the value of weight.
#example for every RM(room) add 3.808865
#A coeff is the value if increased by exactly 1
#the table makes all values weighted the same
#shows the marginal impact of each variable to effect y.*

Out[27]:

	Variable	Weight
5	RM	3.809865
3	CHAS	2.686734
8	RAD	0.306049
1	ZN	0.046420
2	INDUS	0.020559
11	B	0.009312
6	AGE	0.000692
9	TAX	-0.012335
0	CRIM	-0.108011
12	LSTAT	-0.524758
10	PTRATIO	-0.952747
7	DIS	-1.475567
4	NOX	-17.766611

In [28]: *#STANDARDIZING NOTES: Critical step for setting up many models correctly
#Gives all numeric variables a mean of 0, and a variance of 1
#Puts all weights on an equal footing
#Do this by subtracting each value by its mean & divide it by its standard deviation.

#This is the manual way to do it. Use Scalers to preprocess this: See bonus info up tomorrow.*

In [31]: *(X - X.mean()).describe() #Called centering the data.*

Out[31]:

	CRIM	ZN	INDUS	CHAS	NOX	RM
count	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+0
mean	-3.024370e-15	2.076161e-14	-2.800395e-14	-1.189760e-16	2.571505e-16	-8.999389e-1
std	8.601545e+00	2.332245e+01	6.860353e+00	2.539940e-01	1.158777e-01	7.026171e-0
min	-3.607204e+00	-1.136364e+01	-1.067678e+01	-6.916996e-02	-1.696951e-01	-2.723634e+0
25%	-3.531479e+00	-1.136364e+01	-5.946779e+00	-6.916996e-02	-1.056951e-01	-3.991344e-0
50%	-3.357014e+00	-1.136364e+01	-1.446779e+00	-6.916996e-02	-1.669506e-02	-7.613439e-0
75%	6.355894e-02	1.136364e+00	6.963221e+00	-6.916996e-02	6.930494e-02	3.388656e-0
max	8.536268e+01	8.863636e+01	1.660322e+01	9.308300e-01	3.163049e-01	2.495366e+0

In [33]: *X_std = (X - X.mean()) / X.std() #this is how your data should look before feeding it into model
#hovers around 0 and numeric range is small.*

In [34]: `X_std.describe()` *#std is now 1 across the board, Need to make sure data is all on the same scale*

Out[34]:

	CRIM	ZN	INDUS	CHAS	NOX	RM
count	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02
mean	8.326673e-17	3.466704e-16	-3.016965e-15	3.999875e-16	3.167427e-15	-1.258809e-14
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00
min	-4.193669e-01	-4.872402e-01	-1.556302e+00	-2.723291e-01	-1.464433e+00	-3.876413e+00
25%	-4.105633e-01	-4.872402e-01	-8.668328e-01	-2.723291e-01	-9.121262e-01	-5.680681e-01
50%	-3.902803e-01	-4.872402e-01	-2.108898e-01	-2.723291e-01	-1.440749e-01	-1.083583e-01
75%	7.389247e-03	4.872402e-02	1.014995e+00	-2.723291e-01	5.980871e-01	4.822906e-01
max	9.924110e+00	3.800473e+00	2.420170e+00	3.664771e+00	2.729645e+00	3.551530e+00

In [37]: `lreg.fit(X_std, y)`

Out[37]: `LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)`

In [38]: `lreg.coef_`

Out[38]: `array([-0.92906457, 1.08263896, 0.14103943, 0.68241438, -2.05875361, 2.67687661, 0.01948534, -3.10711605, 2.6648522 , -2.07883689, -2.06264585, 0.85010886, -3.74733185])`

```
In [39]: xcoeffs = pd.DataFrame({
          'Variable': X_std.columns,
          'Weight' : lreg.coef_
        }).sort_values(by='Weight', ascending=False)
xcoeffs
```

Out[39]:

	Variable	Weight
5	RM	2.676877
8	RAD	2.664852
1	ZN	1.082639
11	B	0.850109
3	CHAS	0.682414
2	INDUS	0.141039
6	AGE	0.019485
0	CRIM	-0.929065
4	NOX	-2.058754
10	PTRATIO	-2.062646
9	TAX	-2.078837
7	DIS	-3.107116
12	LSTAT	-3.747332

```
In [40]: lreg.score(X_std, y)
#if your r2 did not change your predictions would be practically the same
#the range is reduce TAX and STAT are now larger than NOX
#this is a better way to read and see the coeffs.
#Remeber and increase of 1 is an increase of 1 standard d.
#if you dont take this step (Standardizing) your results will not be valid.

#note this is not to be confused with normalizing, they are different.
```

Out[40]: 0.7406426641094095

```
In [ ]: #When does standardization apply:
         #Linear models
         #Anything that uses a weight penalty (l1, l2, etc)
         #Any algorithm that uses gradient descent
```

```
In [ ]: #Cross Validation Notes:
#How do we know our model will generalize to the outside world?
#We don't know how well our information will work with unknown data.
#Never fit your model to all of your data. #use cross validation to validate your data.
```

```
In [ ]: #Bias vs. Variance Tradeoff: looking for the optimum model
#Bias refer to when your data doesnt adequately capture the info necessary to solve our problem
#Variance refer to when your model mistakenly interprets random correlations as meaningful
#(r2 value looks high but when new data comes in it crashes)
```

Basic idea of Cross Validation: Seperate your data into two groups a test set and training set. Test set is not touched until the very end, and is only used for final model eval Training set is used for fine-tuning and eval your model Note: either randomized or based on date.

```
In [41]: from sklearn.model_selection import train_test_split
```

```
In [45]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=2019)
```

```
In [43]: #train and test are used to split the data where the test size is 20%
#the random state ensure that our random numbers are generated the same
#(randomly shuffled the same way for everyone)
X_train.head()
```

Out[43]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	L
249	0.19073	22.0	5.86	0	0.431	6.718	17.5	7.8265	7	330	19.1	393.74	
51	0.04337	21.0	5.64	0	0.439	6.115	63.0	6.8147	4	243	16.8	393.97	
151	1.49632	0.0	19.58	0	0.871	5.404	100.0	1.5916	5	403	14.7	341.60	
486	5.69175	0.0	18.10	0	0.583	6.114	79.8	3.5459	24	666	20.2	392.68	
235	0.33045	0.0	6.20	0	0.507	6.086	61.5	3.6519	8	307	17.4	376.75	

```
In [46]: lreg.fit(X_train, y_train) #fit on the training sets
```

Out[46]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

```
In [47]: lreg.score(X_test, y_test) #score on the test sets
```

Out[47]: 0.6174065999127849


```
In [ ]: #additional step is to create a validation set: test set within the test set.  
#Think of it as a dress rehearsal for the real test set. You would score on the validation set  
#remember you may not always have the test set (Because this would come later)  
#and this is where you would use validation set
```