



Project 2: Analyzing IMDb Data

Author: Kevin Markham (DC)

For project two, you will complete a series of exercises exploring movie rating data from IMDb.

For these exercises, you will be conducting basic exploratory data analysis on IMDb's movie data, looking to answer such questions as:

What is the average rating per genre? How many different actors are in a movie?

This process will help you practice your data analysis skills while becoming comfortable with Pandas.

Basic level

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

%matplotlib inline
```

Read in 'imdb_1000.csv' and store it in a DataFrame named movies.

```
In [3]: movies = pd.read_csv('/Users/swllms/DAT-10-14-SW/projects/required/project-eda-options/data/imdb_1000.csv')
movies.head()
```

Out[3]:

| | star_rating | title | content_rating | genre | duration | actors_list |
|---|-------------|--------------------------|----------------|--------|----------|---|
| 0 | 9.3 | The Shawshank Redemption | R | Crime | 142 | [u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt... |
| 1 | 9.2 | The Godfather | R | Crime | 175 | [u'Marlon Brando', u'Al Pacino', u'James Caan'] |
| 2 | 9.1 | The Godfather: Part II | R | Crime | 200 | [u'Al Pacino', u'Robert De Niro', u'Robert Duv... |
| 3 | 9.0 | The Dark Knight | PG-13 | Action | 152 | [u'Christian Bale', u'Heath Ledger', u'Aaron E... |
| 4 | 8.9 | Pulp Fiction | R | Crime | 154 | [u'John Travolta', u'Uma Thurman', u'Samuel L.... |

Check the number of rows and columns.

```
In [4]: # Answer:
movies.shape
```

Out[4]: (979, 6)

Check the data type of each column.

```
In [5]: # Answer:
movies.dtypes
```

```
Out[5]: star_rating    float64
title                object
content_rating       object
genre                object
duration             int64
actors_list          object
dtype: object
```

Calculate the average movie duration.

```
In [6]: # Answer:
movies['duration'].mean()
```

```
Out[6]: 120.97957099080695
```

Sort the DataFrame by duration to find the shortest and longest movies.

```
In [7]: # Answer: could have also used nsmallest and nlargest
movies['duration'].nsmallest(1)
```

```
Out[7]: 389      64
Name: duration, dtype: int64
```


```
In [8]: movies['duration'].nlargest(1)
```

```
Out[8]: 476      242
Name: duration, dtype: int64
```

Create a histogram of duration, choosing an "appropriate" number of bins.

```
In [11]: # Answer:
movies['duration'].plot(kind='Hist', bins=30)
```


```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x1a19f3f7f0>
```



Use a box plot to display that same data.

```
In [12]: # Answer:
movies['duration'].plot(kind='box')
```

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x1a19f6cfd0>
```



Intermediate level

Count how many movies have each of the content ratings.


```
In [13]: # Answer:
movies['content_rating'].value_counts()
```

```
Out[13]: R                460
PG-13             189
PG                123
NOT RATED         65
APPROVED          47
UNRATED           38
G                 32
PASSED            7
NC-17             7
X                 4
GP                3
TV-MA             1
Name: content_rating, dtype: int64
```

Use a visualization to display that same data, including a title and x and y labels.

```
In [14]: # Answer:
movies.content_rating.value_counts().plot(kind='barh', title='Content
Rating Counts')
plt.xlabel('Content Rating')
plt.ylabel('Count')
```

```
Out[14]: Text(0, 0.5, 'Count')
```



Convert the following content ratings to "UNRATED": NOT RATED, APPROVED, PASSED, GP.

```
In [15]: # Answer: Condition and result logic
movies['content_rating'] = movies['content_rating'].replace({'NOT RATE
D': 'UNRATED', 'APPROVED': 'UNRATED', 'PASSED': 'UNRATED', 'GP': 'UNRATED'
})
```

```
In [16]: movies['content_rating'].value_counts()
```

```
Out[16]: R          460
PG-13       189
UNRATED     160
PG          123
G           32
NC-17        7
X            4
TV-MA        1
Name: content_rating, dtype: int64
```

Convert the following content ratings to "NC-17": X, TV-MA.

```
In [17]: # Answer: Condition and result logic
movies['content_rating'] = movies['content_rating'].replace({'X': 'NC-17', 'TV-MA': 'NC-17'})
```

```
In [18]: movies['content_rating'].value_counts()
```

```
Out[18]: R          460
PG-13      189
UNRATED    160
PG          123
G           32
NC-17       12
Name: content_rating, dtype: int64
```

Count the number of missing values in each column.

```
In [19]: # Answer:
movies.isnull().sum()
```

```
Out[19]: star_rating      0
title                    0
content_rating           3
genre                    0
duration                 0
actors_list              0
dtype: int64
```

If there are missing values: examine them, then fill them in with "reasonable" values.

```
In [20]: # Answer: 3 missing values;
movies.content_rating.fillna(value='UNRATED', inplace=True)
```

```
In [21]: movies.isnull().sum()
```

```
Out[21]: star_rating      0
         title            0
         content_rating    0
         genre            0
         duration          0
         actors_list       0
         dtype: int64
```

Calculate the average star rating for movies 2 hours or longer, and compare that with the average star rating for movies shorter than 2 hours.

```
In [22]: # Answer:
         movies[movies['duration'] >=120]['star_rating'].mean()
```


```
Out[22]: 7.948898678414082
```

```
In [23]: movies[movies['duration'] <120]['star_rating'].mean()
```

```
Out[23]: 7.838666666666657
```

Use a visualization to detect whether there is a relationship between duration and star rating.

```
In [24]: # Answer:
         movies.plot('duration', 'star_rating', kind='scatter');
```



Calculate the average duration for each genre.

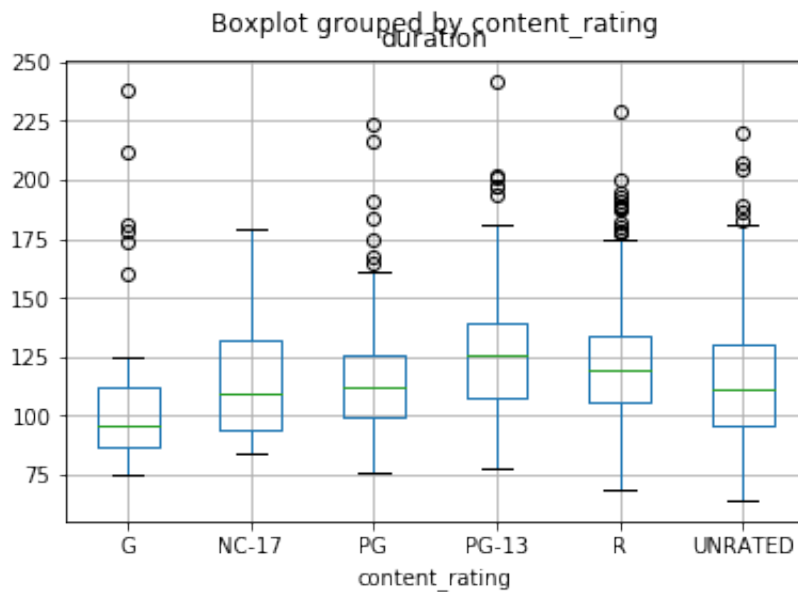
```
In [25]: # Answer: Use groupby for genre and duration and take the mean
movies.groupby('genre')['duration'].mean()
```

```
Out[25]: genre
Action      126.485294
Adventure   134.840000
Animation    96.596774
Biography    131.844156
Comedy       107.602564
Crime        122.298387
Drama        126.539568
Family       107.500000
Fantasy      112.000000
Film-Noir    97.333333
History       66.000000
Horror       102.517241
Mystery      115.625000
Sci-Fi       109.000000
Thriller     114.200000
Western      136.666667
Name: duration, dtype: float64
```

Advanced level

Visualize the relationship between content rating and duration.


```
In [26]: # Answer:
movies.boxplot(column='duration', by='content_rating');
```



Determine the top rated movie (by star rating) for each genre.

```
In [27]: # Answer:
movies.groupby('genre')['title', 'star_rating'].first()
```

Out[27]:

| | title | star_rating |
|-----------|---|-------------|
| genre | | |
| Action | The Dark Knight | 9.0 |
| Adventure | The Lord of the Rings: The Return of the King | 8.9 |
| Animation | Spirited Away | 8.6 |
| Biography | Schindler's List | 8.9 |
| Comedy | Life Is Beautiful | 8.6 |
| Crime | The Shawshank Redemption | 9.3 |
| Drama | 12 Angry Men | 8.9 |
| Family | E.T. the Extra-Terrestrial | 7.9 |
| Fantasy | The City of Lost Children | 7.7 |
| Film-Noir | The Third Man | 8.3 |
| History | Battleship Potemkin | 8.0 |
| Horror | Psycho | 8.6 |
| Mystery | Rear Window | 8.6 |
| Sci-Fi | Blade Runner | 8.2 |
| Thriller | Shadow of a Doubt | 8.0 |
| Western | The Good, the Bad and the Ugly | 8.9 |

Check if there are multiple movies with the same title, and if so, determine if they are actually duplicates.

```
In [28]: # Answer: same title but different actors and movie durations.
movies.duplicated(['title']).sum()
```

Out[28]: 4

```
In [29]: dups = movies[movies.title.duplicated()].title
dups
```

```
Out[29]: 482    The Girl with the Dragon Tattoo
          905                Dracula
          924    Les Misérables
          936    True Grit
Name: title, dtype: object
```

In [30]: `movies[movies.title.isin(dups)]#different duration and different actor list.`

Out[30]:

| | star_rating | title | content_rating | genre | duration | actors_list |
|-----|-------------|---------------------------------|----------------|-----------|----------|--|
| 466 | 7.9 | The Girl with the Dragon Tattoo | R | Crime | 158 | [u'Daniel Craig', u'Rooney Mara', u'Christophe...] |
| 482 | 7.8 | The Girl with the Dragon Tattoo | R | Crime | 152 | [u'Michael Nyqvist', u'Noomi Rapace', u'Ewa Fr...] |
| 662 | 7.7 | True Grit | PG-13 | Adventure | 110 | [u'Jeff Bridges', u'Matt Damon', u'Hailee Stei...] |
| 678 | 7.7 | Les Miserables | PG-13 | Drama | 158 | [u'Hugh Jackman', u'Russell Crowe', u'Anne Hat...] |
| 703 | 7.6 | Dracula | UNRATED | Horror | 85 | [u'Bela Lugosi', u'Helen Chandler', u'David Ma...] |
| 905 | 7.5 | Dracula | R | Horror | 128 | [u'Gary Oldman', u'Winona Ryder', u'Anthony Ho...] |
| 924 | 7.5 | Les Miserables | PG-13 | Crime | 134 | [u'Liam Neeson', u'Geoffrey Rush', u'Uma Thurm...] |
| 936 | 7.4 | True Grit | UNRATED | Adventure | 128 | [u'John Wayne', u'Kim Darby', u'Glen Campbell'] |

Calculate the average star rating for each genre, but only include genres with at least 10 movies

Option 1: manually create a list of relevant genres, then filter using that list

In [9]: `# Answer:`

Option 2: automatically create a list of relevant genres by saving the value_counts and then filtering

```
In [31]: # Answer:
rel_genres = movies['genre'].value_counts()
rel_genres
```

```
Out[31]: Drama          278
Comedy          156
Action          136
Crime           124
Biography        77
Adventure        75
Animation        62
Horror           29
Mystery          16
Western           9
Thriller          5
Sci-Fi           5
Film-Noir         3
Family            2
Fantasy           1
History           1
Name: genre, dtype: int64
```

Option 3: calculate the average star rating for all genres, then filter using a boolean Series

```
In [32]: # Answer:
movies.groupby('genre').star_rating.mean()[movies.genre.value_counts()
>= 10]
```

```
Out[32]: genre
Action          7.884559
Adventure        7.933333
Animation        7.914516
Biography        7.862338
Comedy           7.822436
Crime            7.916935
Drama            7.902518
Horror           7.806897
Mystery          7.975000
Name: star_rating, dtype: float64
```

Option 4: aggregate by count and mean, then filter using the count

```
In [38]: # Answer:
grs = movies.groupby('genre')['star_rating'].agg(['count', 'mean'])
grs[grs['count'] >= 10]
```

Out[38]:

| | count | mean |
|-----------|-------|----------|
| genre | | |
| Action | 136 | 7.884559 |
| Adventure | 75 | 7.933333 |
| Animation | 62 | 7.914516 |
| Biography | 77 | 7.862338 |
| Comedy | 156 | 7.822436 |
| Crime | 124 | 7.916935 |
| Drama | 278 | 7.902518 |
| Horror | 29 | 7.806897 |
| Mystery | 16 | 7.975000 |

Bonus

Figure out something "interesting" using the actors data!

```
In [45]: movies.actors_list.value_counts() ##Order counts and may not count mov
ies with the same cast
###Harry Potter Cast had 6 Movies together, followed by Star Wars, Bac
k to the Future and Toy Stroy
```

```
Out[45]: [u'Daniel Radcliffe', u'Emma Watson', u'Rupert Grint']
6
[u'Mark Hamill', u'Harrison Ford', u'Carrie Fisher']
3
[u'Michael J. Fox', u'Christopher Lloyd', u'Lea Thompson']
2
[u'Tom Hanks', u'Tim Allen', u'Joan Cusack']
2
[u'Ian McKellen', u'Martin Freeman', u'Richard Armitage']
2
[u'Robert Downey Jr.', u'Jude Law', u'Rachel McAdams']
1
[u'Heather Langenkamp', u'Johnny Depp', u'Robert Englund']
```

```

1
[u'Matthew McConaughey', u'Tye Sheridan', u'Jacob Lofland']
1
[u'Ryan Potter', u'Scott Adsit', u'Jamie Chung']
1
[u'Daniel Day-Lewis', u'Madeleine Stowe', u'Russell Means']
1
[u'Gregory Peck', u'Lee Remick', u'Harvey Stephens']
1
[u'Shah Rukh Khan', u'Kajol', u'Amrish Puri']
1
[u'Will Ferrell', u'Emma Thompson', u'Dustin Hoffman']
1
[u'Tom Hanks', u'Tyler Hoechlin', u'Rob Maxey']
1
[u'Patrick Stewart', u'Hugh Jackman', u'Ian McKellen']
1
[u'Richard E. Grant', u'Paul McGann', u'Richard Griffiths']
1
[u'Brad Davis', u'Irene Miracle', u'Bo Hopkins']
1
[u'Patrick Stewart', u'Ian McKellen', u'Hugh Jackman']
1
[u'Catalina Sandino Moreno', u'Guilied Lopez', u'Orlando Tob\x3f3n']
1
[u'Edward Norton', u'Edward Furlong', u"Beverly D'Angelo"]
1
[u'Vladimir Garin', u'Ivan Dobronravov', u'Konstantin Lavronenko']
1
[u'Brad Pitt', u'Robin Wright', u'Jonah Hill']
1
[u'Ulrich M\xfcche', u'Martina Gedeck', u'Sebastian Koch']
1
[u'Sterling Hayden', u'Coleen Gray', u'Vince Edwards']
1
[u'Robert De Niro', u'Charles Grodin', u'Yaphet Kotto']
1
[u'William Holden', u'Gloria Swanson', u'Erich von Stroheim']
1
[u'Jake Gyllenhaal', u'Chris Cooper', u'Laura Dern']
1
[u'Jon Voight', u'Burt Reynolds', u'Ned Beatty']
1
[u'Mel Gibson', u'Danny Glover', u'Gary Busey']
1
[u'Harrison Ford', u'Tommy Lee Jones', u'Sela Ward']
1
..
[u'Sh\xfbichir\xfc4 Moriyama', u'Tokiko Kat\xfc4', u'Sanshi Katsura']

```

```

1
[u'Hayden Christensen', u'Natalie Portman', u'Ewan McGregor']
1
[u'Pam Grier', u'Samuel L. Jackson', u'Robert Forster']
1
[u'Ryan Gosling', u'Michelle Williams', u'John Doman']
1
[u'Paul Giamatti', u'Shari Springer Berman', u'Harvey Pekar']
1
[u'Christian Bale', u'Jennifer Jason Leigh', u'Aitana S\xel\nchez-Gij
\x3n']    1
[u'Susan Sarandon', u'Sean Penn', u'Robert Prosky']
1
[u'Liza Minnelli', u'Michael York', u'Helmut Griem']
1
[u'Joe Pesci', u'Marisa Tomei', u'Ralph Macchio']
1
[u'Tatsuya Fujiwara', u'Aki Maeda', u'Tar\x4 Yamamoto']
1
[u'Graham Chapman', u'John Cleese', u'Eric Idle']
1
[u'Logan Lerman', u'Emma Watson', u'Ezra Miller']
1
[u'Tom Hardy', u'Nick Nolte', u'Joel Edgerton']
1
[u'Paul Newman', u'George Kennedy', u'Strother Martin']
1
[u'Denzel Washington', u'Angela Bassett', u'Delroy Lindo']
1
[u'Mickey Rourke', u'Clive Owen', u'Bruce Willis']
1
[u'Ben Affleck', u'Bryan Cranston', u'John Goodman']
1
[u'Nicole Kidman', u'Paul Bettany', u'Lauren Bacall']
1
[u'Anna Paquin', u'James Van Der Beek', u'Cloris Leachman']
1
[u'Brigitte Helm', u'Alfred Abel', u'Gustav Fr\x6hlich']
1
[u'Ellen Page', u'Michael Cera', u'Jennifer Garner']
1
[u'Sean Connery', u'Michael Caine', u'Christopher Plummer']
1
[u'Nicolas Cage', u'Holly Hunter', u'Trey Wilson']
1
[u'Eric Bana', u'Daniel Craig', u'Marie-Jos\xe9 Croze']
1
[u'Richard Burton', u'Clint Eastwood', u'Mary Ure']
1
[u'Al Pacino', u'Sean Penn', u'Penelope Ann Miller']

```

```

1
[u'Cary Grant', u'Ingrid Bergman', u'Claude Rains']
1
[u'Daniel Day-Lewis', u'Sally Field', u'David Strathairn']
1
[u'Charlton Heston', u'Orson Welles', u'Janet Leigh']
1
[u'Leonardo DiCaprio', u'Djimon Hounsou', u'Jennifer Connelly']
1
Name: actors_list, Length: 969, dtype: int64

```

```
In [55]: hanks = movies.actors_list.str.contains("Tom Hanks")
```

```
In [58]: hanks.value_counts() ###Tom Hanks has been in 14 of the 1000 movies li
sted
```

```
Out[58]: False    965
         True      14
         Name: actors_list, dtype: int64
```

```
In [61]: DH = movies.actors_list.str.contains("Patrick Stewart")
```

```
In [62]: DH.value_counts()
```

```
Out[62]: False    975
         True       4
         Name: actors_list, dtype: int64
```

```
In [14]: movies.star_rating.nsmallest(1)
```

```
Out[14]: 930    7.4
         Name: star_rating, dtype: float64
```

```
In [ ]:
```